

AI-Powered Legal Debate System

1. Approach

The system is designed as an **AI-powered courtroom simulator** where two distinct AI lawyers engage in a structured debate:

- **Prosecution (RAG Lawyer):** Builds **fact-based legal arguments** by retrieving structured information from a small dataset of legal cases (cases.jsonl).
- **Defense (Chaos Lawyer):** Generates **creative and absurd counterarguments**, relying on improvisation rather than evidence.
- **Judge (User + AI Events):** The user plays judge and chooses a winner, while the AI occasionally injects courtroom events (interruptions, warnings) to make the debate lively.

This hybrid approach combines **retrieval-augmented generation (RAG)** with **generative creativity**, ensuring the debate is both informative and entertaining.

2. Retrieval Strategy

The **prosecution lawyer's reasoning** is powered by a **hybrid retrieval strategy** implemented in llm.py:

1. Corpus Expansion:

- The script `generate_facts.py` transforms each legal case into multiple structured factual snippets (year, jurisdiction, evidence snippets).
- These facts are stored in `facts.py` as `DEFAULT_FACTS`.

2. Keyword Filtering:

- Given a case, keywords are extracted and matched against facts.
- This ensures that only facts mentioning relevant terms are considered.

3. Semantic Embedding Search:

- The **SentenceTransformer (MiniLM)** model embeds both case text and candidate facts.
- Cosine similarity is used to rank facts by relevance.

4. Hybrid Ranking:

- Keyword + embedding filtering ensures robustness.
- Top-k facts (e.g., 5–10) are selected as evidence.

5. Fact-based Argument Generation:

- One fact is selected per round and wrapped in a legal argument template:
“Based on evidence, [fact]. This strongly supports the prosecution's case.”

This ensures that prosecution arguments remain grounded in structured case knowledge, unlike the defense.

. Metadata Structure

The **metadata schema** for each case in cases.jsonl is:

```
{  
  "id": 1,  
  "title": "Case Title",  
  "year": "1998",  
  "jurisdiction": "Supreme Court",  
  "tags": ["contract", "liability"],  
  "text": "Full case description..."  
}
```

From this structure, multiple **fact snippets** are generated (e.g., *“Case X was decided in the Supreme Court in 1998”*), which improves retrieval interpretability.

4. System Design

The architecture follows a **modular client–server design**:

- **Backend (FastAPI)**
 - Orchestrates debates (main.py).
 - Loads the corpus (retrieval.py).
 - Generates cases (generator.py).
 - Produces lawyer arguments (llm.py).
 - Summarizes verdicts.
 - Uses **timeout safeguards** so each round completes even if an AI call hangs.
- **Frontend (Streamlit)**
 - Provides interactive UI.
 - Lets users generate cases, start debates, view prosecution vs. defense arguments, and choose verdicts.
 - Displays judge’s random interventions and stores debate history.

- **RAG Lawyer vs. Chaos Lawyer**
 - RAG Lawyer = grounded, evidence-based, relies on hybrid retrieval.
 - Chaos Lawyer = imaginative, absurd, generates entertainment value.
- **Judge**
 - Human judge selects winner.
 - AI judge adds random interventions for realism.

5. Challenges Faced

- **No prior knowledge to Streamlit and Fastapi. Having to learn both frontend and backend was probably the biggest issue for me but pretty useful in the future.**
- **Having to use a mock llm due to memory issues and no gpu. Also didn't have access to api keys of Openai, Anthropic, etc.**
- **Though I had basic understanding of an LLM, this was my first real-time application of it and combining it with the lack of web development knowledge made it pretty tough.**